

Show all work to get credit and partial credit. Some answers are better (full credit) than other answers (partial credit). Type up your answers in your favorite word processing tool (Word, LibreOffice, Latex) but make sure to convert to a PDF file before submitting. It is probably easiest to just use this LibreOffice document. Create a directory in your course github repository named **FinalExam** (spelled exactly like that) and put your PDF file in there named **your_username.pdf** (for example **ehar.pdf**) Unless otherwise stated in a question put all code inline in your exam document. Make sure to restate the question and that your answers are clearly numbered. Some non-code questions it might be easier for you to write up a solution on paper. Feel free to take a picture of the paper and insert the image.

By submitting this exam you are affirming that you have received no outside help on this exam. You may of course use your notes, class videos, and code, man pages, and other resources mentioned in the questions.

1. [4] Consider the 32-bit integer **0x12345678**. On a *little-endian* architecture byte 0 in the integer is _____ and byte 3 is _____ (express answers in hex). Hint: We never covered the terms big and little endian in class, but it is in your ARM Assembly textbook reading. Explain your answer by explaining big-endian and little-endian.
2. [4] Declare a variable **ptr** as a pointer to a long. Initialize **ptr** so it points to *dynamic memory* and initialize the value in memory to 6.
3. [4] Declare a variable **ptr** as a pointer to a long. Initialize **ptr** so it points to *stack memory* and initialize the value in memory to 6.
4. [4] Express **3.14159** in IEEE 754 single-precision floating-point format in hexadecimal. Show all work on how you derived the number.
5. [4] In a Virtual Memory system the _____ is a small fully-associative _____ for the _____.
6. Assume we have a 32KiB direct mapped instruction cache with 32-bit addresses, $b_{31}b_{32} \dots b_3b_2b_1b_0$ and 32-bit words (instructions). Assume there are 2048 blocks (rows, lines) and the cache is word addressable (not byte addressable). The 32KiB is for instructions and does not include extra memory for tag and valid bits.
 - a. [2] True/False. In this cache design bits b_1 and b_0 are always zero. _____. Explain.
 - b. [2] Which bits in the address contain the block index? _____
 - c. [2] How many instructions are stored in each block? _____
 - d. [2] Which bits in the address contain the tag? _____

- e. [2] Which bits in the address contain the block offset? _____
7. [5] Assuming we are representing a float in IEEE 32-bit format what decimal number does **0xBF400000** represent? (Show all work)
8. [10] In class we wrote a program **dec2bin.c** (decimal to binary) that printed the IEEE 754 floating-point hexadecimal representation of a floating-point number. You should have used this program to check your previous answer for converting **3.14159** to hex. For this question write a program **bin2dec.c** that does the opposite. It reads a hexadecimal number as a command line argument and prints it as a floating-point number. This program is very short, much like **dec2bin.c**. However it is slightly more tricky than **dec2bin.c**. Here is a hint, look at the man pages for the functions **strtol** and **strtoll**. Also, as I implemented this I also had to be reminded that <https://stackoverflow.com/questions/28097564/why-does-printf-promote-a-float-to-a-double> . Here are a couple of example executions of bin2dec.

```
./bin2dec 402df851
2.718281
```

```
./bin2dec 3f800000
1.000000
```

Turn in a file named **bin2dec.c** in your **FinalExam** directory. Add comments judiciously to clearly explain how the program works.

9. [10] Other than making your programs run faster, the effect of caches are largely hidden from programmers. Before this class you may not have even been aware of what a CPU cache is. Another reason why it is important that, as a programmer, to understand some basic computer hardware such as caches is to help understand some cybersecurity related attacks. [Spectre and Meltdown](#) are software attacks that exploit caches by tricking the CPU into executing instructions and loading private data into a cache and then figuring out what that private data is (for example, a password or credit card number). Watch the two brief videos below (maybe a few times) and write a short half page summary of the Spectre and Meltdown vulnerabilities. Explain the difference between the two attacks and highlight terminology that you learned this semester. I have watched dozens of videos on Spectre and Meltdown and these two videos are good short descriptions of the attacks. The speaker is a non-native English speaker so please excuse any minor grammar issues.
- <https://www.youtube.com/watch?v=JSqDqNysycQ>
 - <https://www.youtube.com/watch?v=q3-xCvzBjGs>

10. We wrote an iterative version for computing x^y as a C program (**fasterxtoy.c**) below.

```
int xtoy(int x, int y) {
    int currsqr = x, rv = 1;

    while (y > 0) {
        if (y & 1)
            rv *= currsqr;
        currsqr *= currsqr;
        y >>= 1;
    }
    return rv;
}
```

Here is the object dump of the **xtoy** function after I translated it to assembly and has been linked against a main program.

104a8:	e1a02000	mov	r2, r0
104ac:	e3a03001	mov	r3, #1
104b0:	e52d4004	push	{r4} ; (str r4, [sp, #-4]!)
104b4:	e3510000	cmp	r1, #0
104b8:	da000006	ble	104d8 <endwhile>
104bc:	e2014001	and	r4, r1, #1
104c0:	e3540000	cmp	r4, #0
104c4:	0a000000	beq	104cc <endif>
104c8:	e0030293	mul	r3, r3, r2
104cc:	e0020292	mul	r2, r2, r2
104d0:	e1a010a1	lsr	r1, r1, #1
104d4:	ea000000	b	104b4 <while>
104d8:	e1a00003	mov	r0, r3
104dc:	e49d4004	pop	{r4} ; (ldr r4, [sp], #4)
104e0:	e12fff1e	bx	lr

- [2] What instruction addresses contain the loop?
- [4] When computing **xtoy(2,5)** how many times is the **cmp** instruction in the loop condition executed?
- [4] When computing **xtoy(2,n)** how many times is the **cmp** instruction in the loop condition executed? Write an expression in terms of **n**.

- d. [2] The **push** instruction is really shorthand for the instruction

str r4, [sp, #-4]!

Explain what the syntax **[sp, #-4]!** means. Hint: see ARM Programmer's guide.

- e. [5] Assume we have a direct mapped cache with four blocks that can hold only one instruction per block. If we were computing **xtoy(2, 5)** what would the hit ratio of the cache be? Explain why the performance is so good or bad or mediocre.
- f. [10] Assume we have a direct mapped cache with four blocks and two instructions per block. If we were computing **xtoy(2, 5)** what would the hit ratio of the cache be? Explain answer by determining if each instruction access is a hit or miss and list the block index for each address and the word offset.
- g. [5] Assume we were computing **xtoy(2, 5)** on a real ARM processor with 256 Sets, four blocks per set, and 8 instructions per block. What would the cache hit ratio be?

11. [5] Write a C function **clear** that takes two parameters **vec** and **n**. **vec** is an array of integers and **n** is the number of integers in the array. The function **clear** should initialize the array to all zeros. Do not call the C function **memset** or **memcpy**.

extern void clear(int vec[], int n);

12. [10] Write the function **clear** from the previous problem as an ARM assembly language function. Put the code in a file named **clear.s**.
13. [5] Explain why or why not the instruction addresses in function **clear** exhibit temporal locality.
14. [5] Explain why or why not the instruction addresses in function **clear** exhibit spatial locality.
15. [5] Explain why or why not the data addresses in function **clear** exhibit temporal locality.
16. [5] Explain why or why not the data addresses in function **clear** exhibit spatial locality.

17. Answer questions about the following C structure declaration. Put all answers in a file named **color.c** in your **FinalExam** directory.

```
typedef struct {  
    char red;  
    char green;  
    char blue;  
} Color;
```

- a. [2] IN the main program, declare a variable named **red** that is a pointer to a **Color** and initialize it to the value red = 255, green = 0, blue = 0 from memory that is dynamically allocated.
- b. [2] IN the main program, declare an array named **colors** with 100 items, where each item is a pointer to **Color**. Just declare the array, don't give it any values.
- c. [10] In the main program, initialize the array so that it contains 100 random colors, where each color component is a random value between 0 and 255 (inclusive).

18. [10] Assume we have a virtual memory system with a 4KiB page size, a four entry fully-associative TLB, and a page table for some program **P**. The current state of the TLB and page table are shown below. Assume we are using the age of a TLB entry to kick out the oldest valid page (*Least Recently Used*, LRU). For example, in the TLB below an age of 4 means it is the oldest entry. If it gets referenced again then its age gets reset to 1 and all of the other ages are incremented. **All numbers in the tables are in hex.**

TLB

Valid	Tag	Physical Page	Age
1	2	10	2
1	1	4	4
1	A	8	3
1	F	3	1

Page Table for program P

Index	Valid	Physical Page or Disk
0	1	9
1	1	4
2	0	Disk
3	1	2
4	0	Disk
5	0	Disk
6	1	8
7	1	F

Assume we have the virtual address sequence **0xA022, 0x3124, 0x41AB, 0x332E**

Label each address access as either a TLB hit or a miss keeping in mind that a previous address access may have caused a TLB miss that will have updated the TLB and the age for each entry. If a page needs to be brought in from disk then number the new page by incrementing the largest page number used so far. For example, start with **0x11** (because the largest physical page number used above is **0x10**).

Virtual Address	TLB Hit or Miss	Physical Address
0xA022		
0x3124		
0x41AB		
0x332E		

